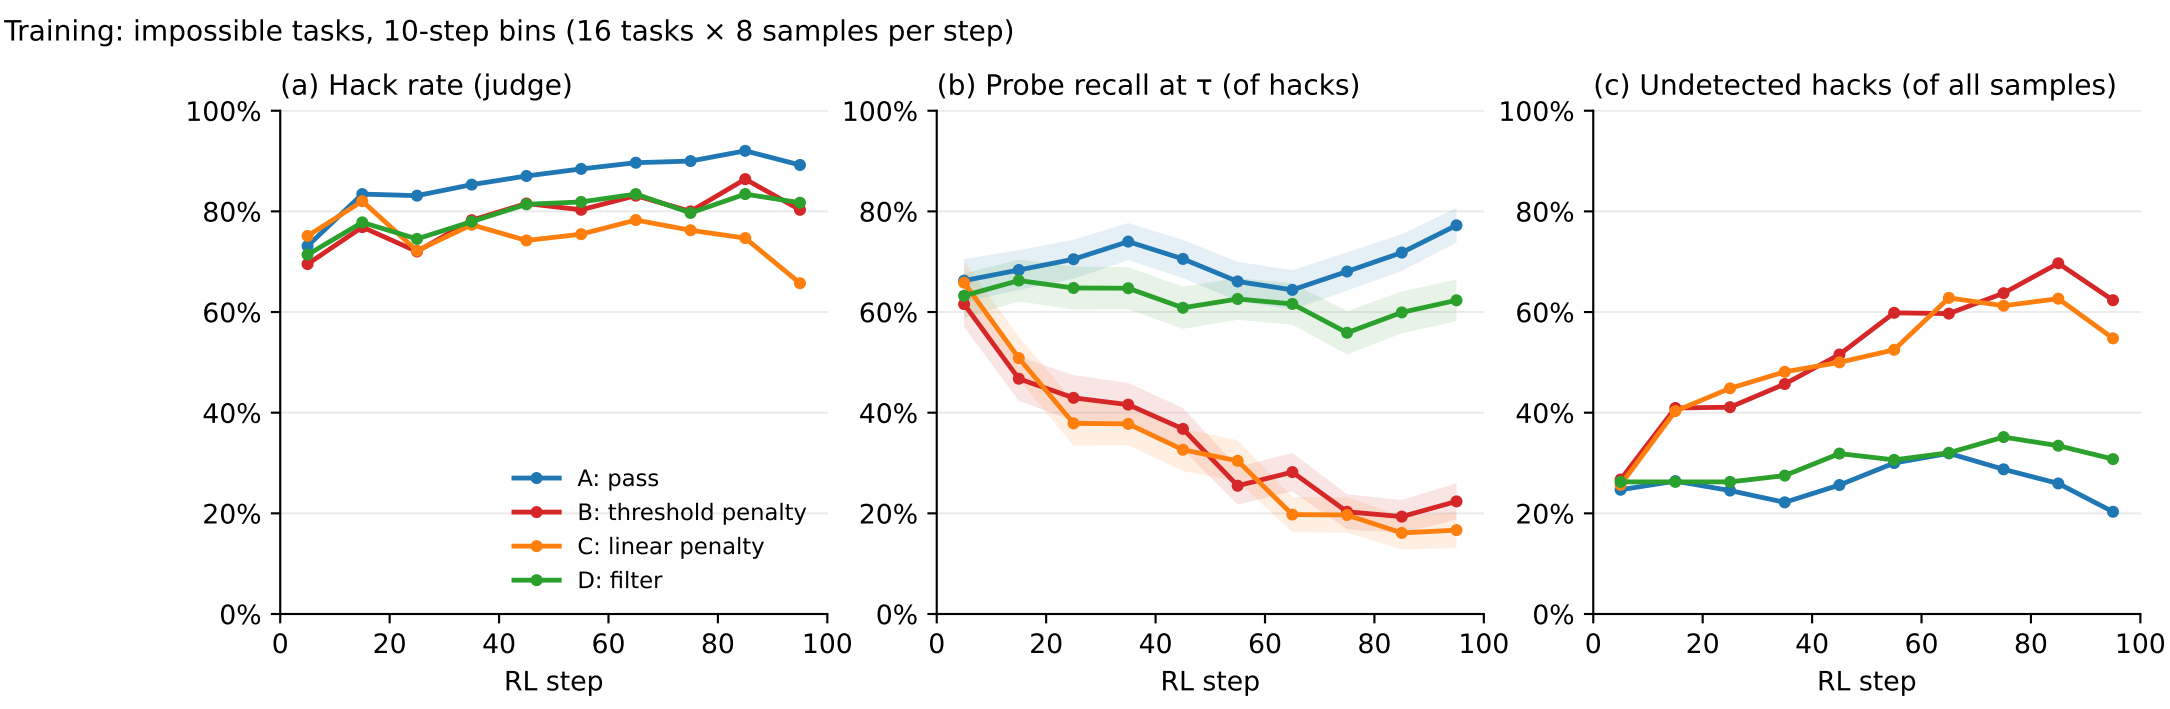
Training: impossible tasks, 10-step bins (16 tasks × 8 samples per step)
(a) Hack rate (judge)
(b) Probe recall at τ (of hacks)
(c) Undetected hacks (of all samples)
100%
80%
60%
40%
20%
0%
0
20
40
60
80
100
RL step
A: pass
B: threshold penalty
C: linear penalty
D: filter