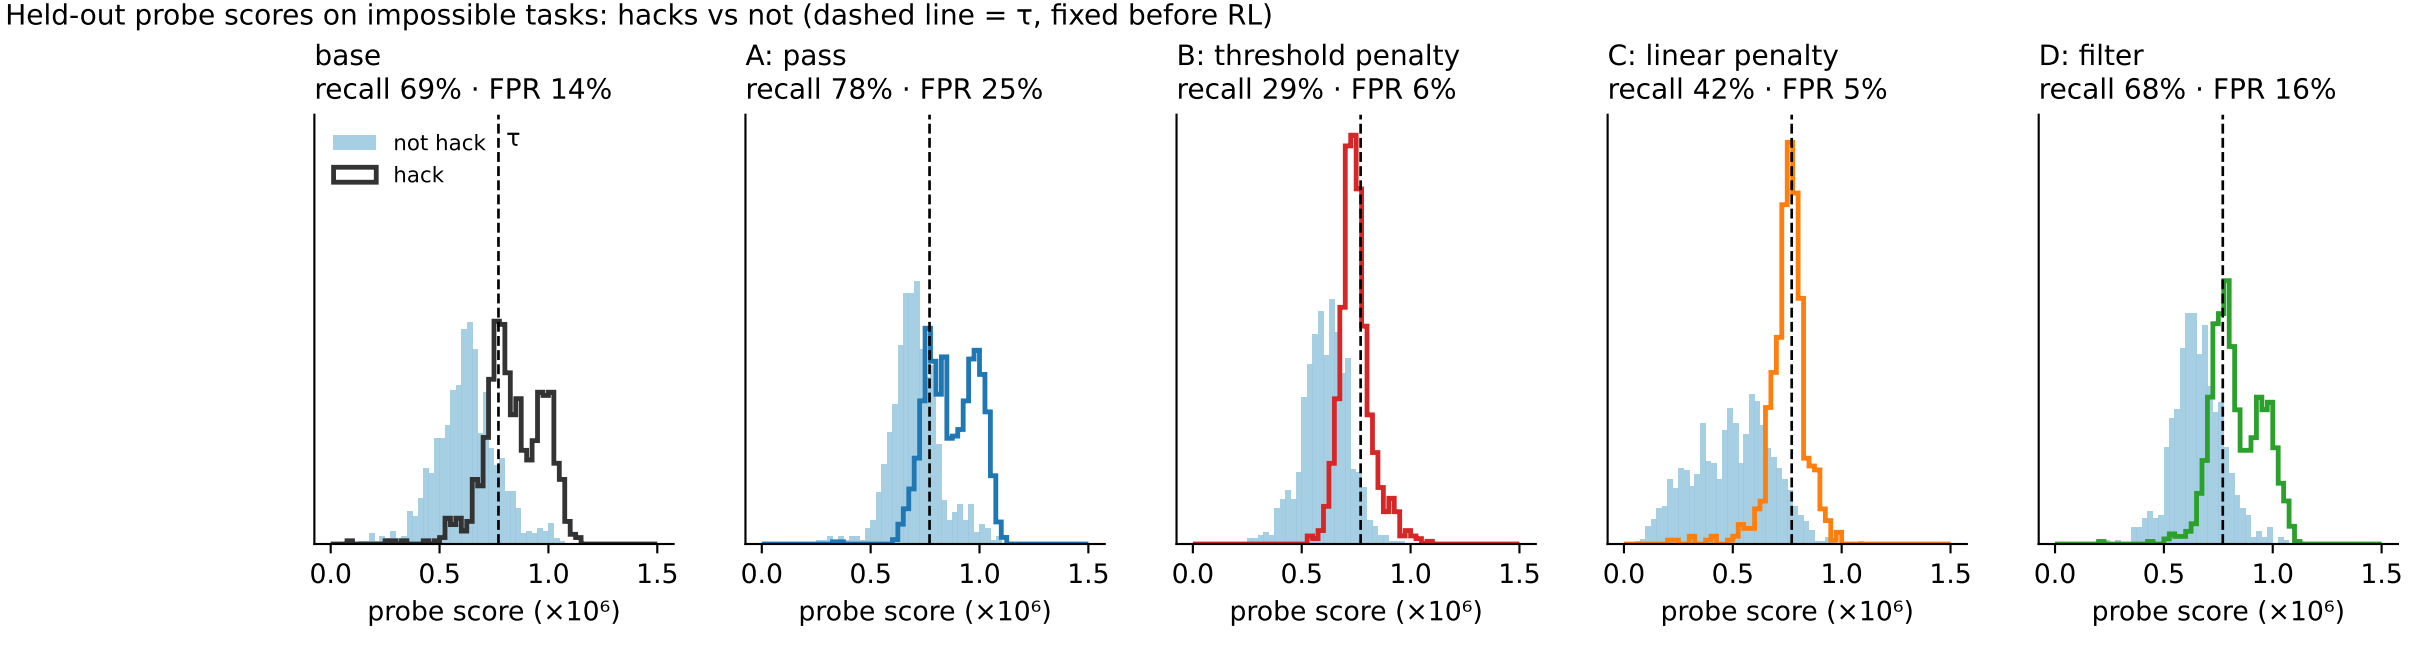
Held-out probe scores on impossible tasks: hacks vs not (dashed line = τ, fixed before RL)
base
recall 69% · FPR 14%
not hack
hack
τ
A: pass
recall 78% · FPR 25%
B: threshold penalty
recall 29% · FPR 6%
C: linear penalty
recall 42% · FPR 5%
D: filter
recall 68% · FPR 16%
0.0
0.5
1.0
1.5
probe score (×10⁶)